

Guobin Shen

Room 208A

Phone: (+86) 139-3142-5808

Automation Mansion

Email: shenguobin2021@ia.ac.cn

No. 95, Zhongguancun East Road, Haidian

Alt: floyed_shen@outlook.com

Beijing, China

Homepage: floyedshen.github.io

RESEARCH INTERESTS

I am passionate about leveraging neuroscience and cognitive science to build scalable, trustworthy, and safe AI systems. My research focuses on developing brain-inspired approaches to understand and improve large models, with emphasis on alignment methods, uncertainty quantification, and robustness against failure modes such as jailbreak attacks and hallucinations. I aim to build AI systems that are not only powerful but also scalable, interpretable, and reliable.

ACADEMIC APPOINTMENTS

Human Intelligence (Hi) Lab, RedNote

Beijing, China

Ace Top Intern Program

September 2021 – Present

Responsibilities: Developed scalable AI alignment methods using efficient human feedback and automated preference learning.

EDUCATION

Institute of Automation, Chinese Academy of Sciences

Beijing, China

Ph. D. in Machine Learning

September 2021 – June 2026 (expected)

Advisor: [Prof. Yi Zeng](#).

Sun Yat-sen University

Guangzhou, China

B. Eng. in Communication Engineering

September 2017 – June 2021















Advisor: [Prof. Xiang Chen](#).

Grade: 1/85 of graduating class.





PUBLICATIONS

LLM Alignment & AI Safety:

1. **Shen, Guobin**, Zhao, Dongcheng, Dong, Yiting, He, Xiang, and Zeng, Yi. “Jailbreak Antidote: Runtime Safety-Utility Balance via Sparse Representation Adjustment in Large Language Models.” *Proceedings of the 13th International Conference on Learning Representations (ICLR)*, 2025. [\[OpenReview\]](#) [\[PDF\]](#)
2. **Shen, Guobin**, Zhao, Dongcheng, He, Xiang, Feng, Linghao, Dong, Yiting, Wang, Jihang, Zhang, Qian, and Zeng, Yi. “Neuro-Vision to Language: Image Reconstruction and





- Interaction via Non-invasive Brain Recordings.” *Proceedings of the 38th Conference on Neural Information Processing Systems (NeurIPS)*, 2024.  [\[PDF\]](#)  [\[Poster\]](#)
3. **Shen, Guobin**, Zhao, Dongcheng, Bao, Aorigele, He, Xiang, Dong, Yiting, and Zeng, Yi. “StressPrompt: Does Stress Impact Large Language Models and Human Performance Similarly?” *Proceedings of the 39th AAAI Conference on Artificial Intelligence (AAAI)*, 2025.  [\[OpenReview\]](#)  [\[PDF\]](#)
 4. **Shen, Guobin**, Zhao, Dongcheng, Dong, Yiting, Zhang, Qian, and Zeng, Yi. “Convergent Evolution across Modalities, Scales and Training Trajectories: Evidence for Human Brain-AI Alignment”, 2025.  [\[Arxiv\]](#)
 5. **Shen, Guobin**, Zhao, Dongcheng, Tong, Haibo, Li, Jindong, Zhao, Feifei, and Zeng, Yi. “Safety Instincts: LLMs Learn to Trust Their Internal Compass for Self-Defense.” *arXiv preprint arXiv:2510.01088*, 2025.  [\[Arxiv\]](#)
 6. **Shen, Guobin**, Zhao, Dongcheng, Feng, Linghao, He, Xiang, Wang, Jihang, Shen, Sicheng, Tong, Haibo, Dong, Yiting, Li, Jindong, Zheng, Xiang, and others. “PandaGuard: Systematic Evaluation of LLM Safety in the Era of Jailbreaking Attacks.” *arXiv preprint arXiv:2505.13862*, 2025.  [\[Project\]](#)  [\[Arxiv\]](#)  [\[Code\]](#)  [\[Dataset\]](#)
 7. **Shen, Guobin**, Zhao, Dongcheng, Dong, Yiting, Li, Yang, Li, Jindong, Sun, Kang, and Zeng, Yi. “Astrocyte-Enabled Advancements in Spiking Neural Networks for Large Language Modeling.” *arXiv preprint arXiv:2312.07625*, 2023.  [\[Arxiv\]](#)
 8. Wu, Ping, **Shen, Guobin**, Zhao, Dongcheng, Wang, Yuwei, Dong, Yiting, Shi, Yu, Lu, Enmeng, Zhao, Feifei, and Zeng, Yi. “CVC: A Large-Scale Chinese Value Rule Corpus for Value Alignment of Large Language Models.” *arXiv preprint arXiv:2506.01495*, 2025.  [\[Arxiv\]](#)  [\[Code\]](#)  [\[Dataset\]](#)

Spiking Neural Networks & Brain-Inspired AI:

9. **Shen, Guobin**, Zhao, Dongcheng, Dong, Yiting, and Zeng, Yi. “Brain-Inspired Neural Circuit Evolution for Spiking Neural Networks.” *Proceedings of the National Academy of Sciences*, vol. 120, no. 39, 2023, p. e2218173120. National Academy of Sciences.  [\[PDF\]](#)
10. **Shen, Guobin**, Zhao, Dongcheng, Li, Tenglong, Li, Jindong, and Zeng, Yi. “Are Conventional SNNs Really Efficient? A Perspective from Network Quantization.” *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, 2024, pp. 27538-27547.  [\[PDF\]](#)  [\[Poster\]](#)
11. **Shen, Guobin**, Zhao, Dongcheng, Dong, Yiting, Li, Yang, Zhao, Feifei, and Zeng, Yi. “Learning the Plasticity: Plasticity-Driven Learning Framework in Spiking Neural Networks.” *Advances in Neural Information Processing Systems (NeurIPS)*, 2025.  [\[PDF\]](#)
12. **Shen, Guobin**, Zhao, Dongcheng, and Zeng, Yi. “Backpropagation with Biologically Plausible Spatiotemporal Adjustment for Training Deep Spiking Neural Networks.” *Patterns*, vol. 3, no. 6, 2022. Elsevier.  [\[PDF\]](#)
13. **Shen, Guobin**, Zhao, Dongcheng, and Zeng, Yi. “Exploiting Nonlinear Dendritic Adaptive Computation in Training Deep Spiking Neural Networks.” *Neural Networks*, vol. 170, 2024, pp. 190-201. Pergamon.  [\[PDF\]](#)
14. **Shen, Guobin**, Zhao, Dongcheng, and Zeng, Yi. “Exploiting High-Performance Spiking Neural Networks with Efficient Spiking Patterns.” *IEEE Transactions on Emerging Topics in Computational Intelligence (TETCI)*, 2025.
15. **Shen, Guobin**, Zhao, Dongcheng, Shen, Sicheng, and Zeng, Yi. “Enhancing Spiking Transformers with Binary Attention Mechanisms.” *The Second Tiny Papers Track at ICLR 2024*.  [\[PDF\]](#)
16. **Shen, Guobin**, Zhao, Dongcheng, Dong, Yiting, Li, Yang, and Zeng, Yi. “Dive into the Power of Neuronal Heterogeneity.” *arXiv preprint arXiv:2305.11484*, 2023.  [\[Arxiv\]](#)




17. Zhao, Dongcheng, **Shen, Guobin**, Dong, Yiting, Li, Yang, and Zeng, Yi. “Improving Stability and Performance of Spiking Neural Networks through Enhancing Temporal Consistency.” *Pattern Recognition*, vol. 159, 2025, p. 111094. Pergamon.  [\[Arxiv\]](#)  [\[PDF\]](#)
18. Han, Bing, Zhao, Feifei, Zeng, Yi, and **Guobin Shen**. “Developmental Plasticity-Inspired Adaptive Pruning for Deep Spiking and Artificial Neural Networks.” *IEEE Transactions on Pattern Analysis and Machine Intelligence*, 2024. IEEE.  [\[PDF\]](#)
19. Pan, Wenxuan, Zhao, Feifei, **Shen, Guobin**, Han, Bing, and Zeng, Yi. “Brain-Inspired Multi-Scale Evolutionary Neural Architecture Search for Deep Spiking Neural Networks.” *IEEE Transactions on Evolutionary Computation*, 2024. IEEE.
20. Shen, Sicheng, Zhao, Dongcheng, **Shen, Guobin**, and Zeng, Yi. “TIM: An Efficient Temporal Interaction Module for Spiking Transformer.” *Proceedings of the 33rd International Joint Conference on Artificial Intelligence (IJCAI 2024)*, 2024.  [\[PDF\]](#)
21. He, Xiang, Liu, Xiangxi, Li, Yang, Zhao, Dongcheng, **Shen, Guobin**, Kong, Qingqun, Yang, Xin, and Zeng, Yi. “CACE-Net: Co-guidance Attention and Contrastive Enhancement for Effective Audio-Visual Event Localization.” *Proceedings of the 32nd ACM International Conference on Multimedia (MM)*, 2024, pp. 985-993.  [\[OpenReview\]](#)  [\[PDF\]](#)
22. He, Xiang, Zhao, Dongcheng, Li, Yang, **Shen, Guobin**, Kong, Qingqun, and Zeng, Yi. “An Efficient Knowledge Transfer Strategy for Spiking Neural Networks from Static to Event Domain.” *Proceedings of the AAAI Conference on Artificial Intelligence (AAAI)*, vol. 38, no. 1, 2024, pp. 512-520.  [\[Arxiv\]](#)
23. Han, Bing, Zhao, Feifei, Zeng, Yi, Pan, Wenxuan, and **Shen, Guobin**. “Enhancing Efficient Continual Learning with Dynamic Structure Development of Spiking Neural Networks.” *Proceedings of the 32nd International Joint Conference on Artificial Intelligence (IJCAI)*, 2023.  [\[PDF\]](#)
24. Zeng, Yi, Zhao, Dongcheng, Zhao, Feifei, **Shen, Guobin**, Dong, Yiting, Lu, Enmeng, Zhang, Qian, Sun, Yinqian, Liang, Qian, Zhao, Yuxuan, and others. “BrainCog: A Spiking Neural Network Based, Brain-Inspired Cognitive Intelligence Engine for Brain-Inspired AI and Brain Simulation.” *Patterns*, 2023, p. 100789.  [\[PDF\]](#)
25. Shen, Sicheng, Zhao, Dongcheng, Feng, Linghao, Yue, Zeyang, Li, Jindong, Li, Tenglong, **Shen, Guobin**, and Zeng, Yi. “STEP: A Unified Spiking Transformer Evaluation Platform for Fair and Reproducible Benchmarking.” *Advances in Neural Information Processing Systems (NeurIPS) Dataset and Benchmark Track*, 2025.  [\[PDF\]](#)

Hardware Acceleration & System Optimization:

26. **Shen, Guobin**, Li, Jindong, Li, Tenglong, Zhao, Dongcheng, and Zeng, Yi. “SpikePack: Enhanced Information Flow in Spiking Neural Networks with High Hardware Compatibility.” *Proceedings of the IEEE/CVF International Conference on Computer Vision (ICCV)*, 2025.  [\[PDF\]](#)
27. Li, Jindong, **Shen, Guobin**, Zhao, Dongcheng, Zhang, Qian, and Zeng, Yi. “Firefly: A High-Throughput Hardware Accelerator for Spiking Neural Networks with Efficient DSP and Memory Optimization.” *IEEE Transactions on Very Large Scale Integration (VLSI) Systems*, vol. 31, no. 8, 2023, pp. 1178-1191. IEEE.  [\[PDF\]](#)
28. Li, Jindong, **Shen, Guobin**, Zhao, Dongcheng, Zhang, Qian, and Zeng, Yi. “Firefly v2: Advancing Hardware Support for High-Performance Spiking Neural Network with a Spatiotemporal FPGA Accelerator.” *IEEE Transactions on Computer-Aided Design of Integrated Circuits and Systems*, 2024. IEEE.  [\[PDF\]](#)
29. Li, Tenglong, Li, Jindong, **Shen, Guobin**, Zhao, Dongcheng, Zhang, Qian, and Zeng, Yi. “FireFly-S: Exploiting Dual-Side Sparsity for Spiking Neural Networks Acceleration with Reconfigurable Spatial Architecture.” *IEEE Transactions on Circuits and Systems I: Regular Papers*, 2024. IEEE.  [\[PDF\]](#)

30. Li, Jindong, Li, Tenglong, **Shen, Guobin**, Zhao, Dongcheng, Zhang, Qian, and Zeng, Yi. "Revealing Untapped DSP Optimization Potentials for FPGA-Based Systolic Matrix Engines." *2024 34th International Conference on Field-Programmable Logic and Applications (FPL)*, IEEE, 2024, pp. 197-203.  [\[Arxiv\]](#)  [\[PDF\]](#)
31. Li, Jindong, Li, Tenglong, **Shen, Guobin**, Zhao, Dongcheng, Zhang, Qian, and Zeng, Yi. "Pushing Up to the Limit of Memory Bandwidth and Capacity Utilization for Efficient LLM Decoding on Embedded FPGA." *2025 Design, Automation & Test in Europe Conference (DATE)*, IEEE, 2025, pp. 1-7.  [\[PDF\]](#)
32. Li, Jindong, Li, Tenglong, Chen, Ruiqi, **Shen, Guobin**, Zhao, Dongcheng, Zhang, Qian, and Zeng, Yi. "Hummingbird: A Smaller and Faster Large Language Model Accelerator on Embedded FPGA." *The 2025 International Conference on Computer-Aided Design (ICCAD)*, 2025.  [\[PDF\]](#)

Datasets & Data Augmentation:

33. **Shen, Guobin**, Zhao, Dongcheng, and Zeng, Yi. "EventMix: An Efficient Data Augmentation Strategy for Event-Based Learning." *Information Sciences*, vol. 644, 2023, p. 119170. Elsevier.  [\[PDF\]](#)
34. Dong, Yiting, He, Xiang, **Shen, Guobin**, Zhao, Dongcheng, Li, Yang, and Zeng, Yi. "EventZoom: A Progressive Approach to Event-Based Data Augmentation for Enhanced Neuromorphic Vision." *Proceedings of the 39th AAAI Conference on Artificial Intelligence (AAAI)*, 2025.  [\[OpenReview\]](#)
35. Dong, Yiting, Li, Yang, Zhao, Dongcheng, **Shen, Guobin**, and Zeng, Yi. "Bullying10K: A Large-Scale Neuromorphic Dataset Towards Privacy-Preserving Bullying Recognition." *Advances in Neural Information Processing Systems (NeurIPS)*, vol. 36, 2024.  [\[PDF\]](#)

ACADEMIC SERVICES

Serve as a reviewer for conferences including *NeurIPS*, *ICML*, *ICLR*, *CVPR*, *ICCV*, *ECCV*, *AAAI*, *MM*, *AISTATS*, and journals including *IEEE Computational Intelligence Magazine*, *Pattern Recognition*, *Neural Networks*, and *Neurocomputing*.

TEACHING

University of Chinese Academy of Sciences

July 2023 – December 2023

Teaching Assistant, Systems and Computational Neuroscience

AWARDS AND HONORS

- **Best Paper Award for Chinese Scientists, *Cell Press* (2022)**
- **Best Paper Award, *Cell Press* (2023)**
- **Chinese Academy of Sciences President Scholarship (2025)**
 - Academic honor from Chinese Academy of Sciences, recognizing doctoral students with outstanding academic achievements (top 1%)
- **National Scholarship (Doctoral Student) (2024)**
 - Granted for exceptional research contributions and academic excellence (top 1%).
- **National Scholarship (Undergraduate) (2019, 2020)**
 - Awarded by the Chinese Government for outstanding performance in academics, extracurriculars, and leadership (top 2%).
- **National Second Prize, National Undergraduate Electronic Design Competition (2019)**
- **Runner-Up, International Aerial Robotics Competition (2019)**